Learning curves, random selection (mean over seeds, band = 95% bootstrap over 70 seeds per block).
Block N: seeds 3000-3034 + 5000-5034 (default head, aw8, aw8_all); block S: seeds 8300-8334 + 8350-8384 (all others).

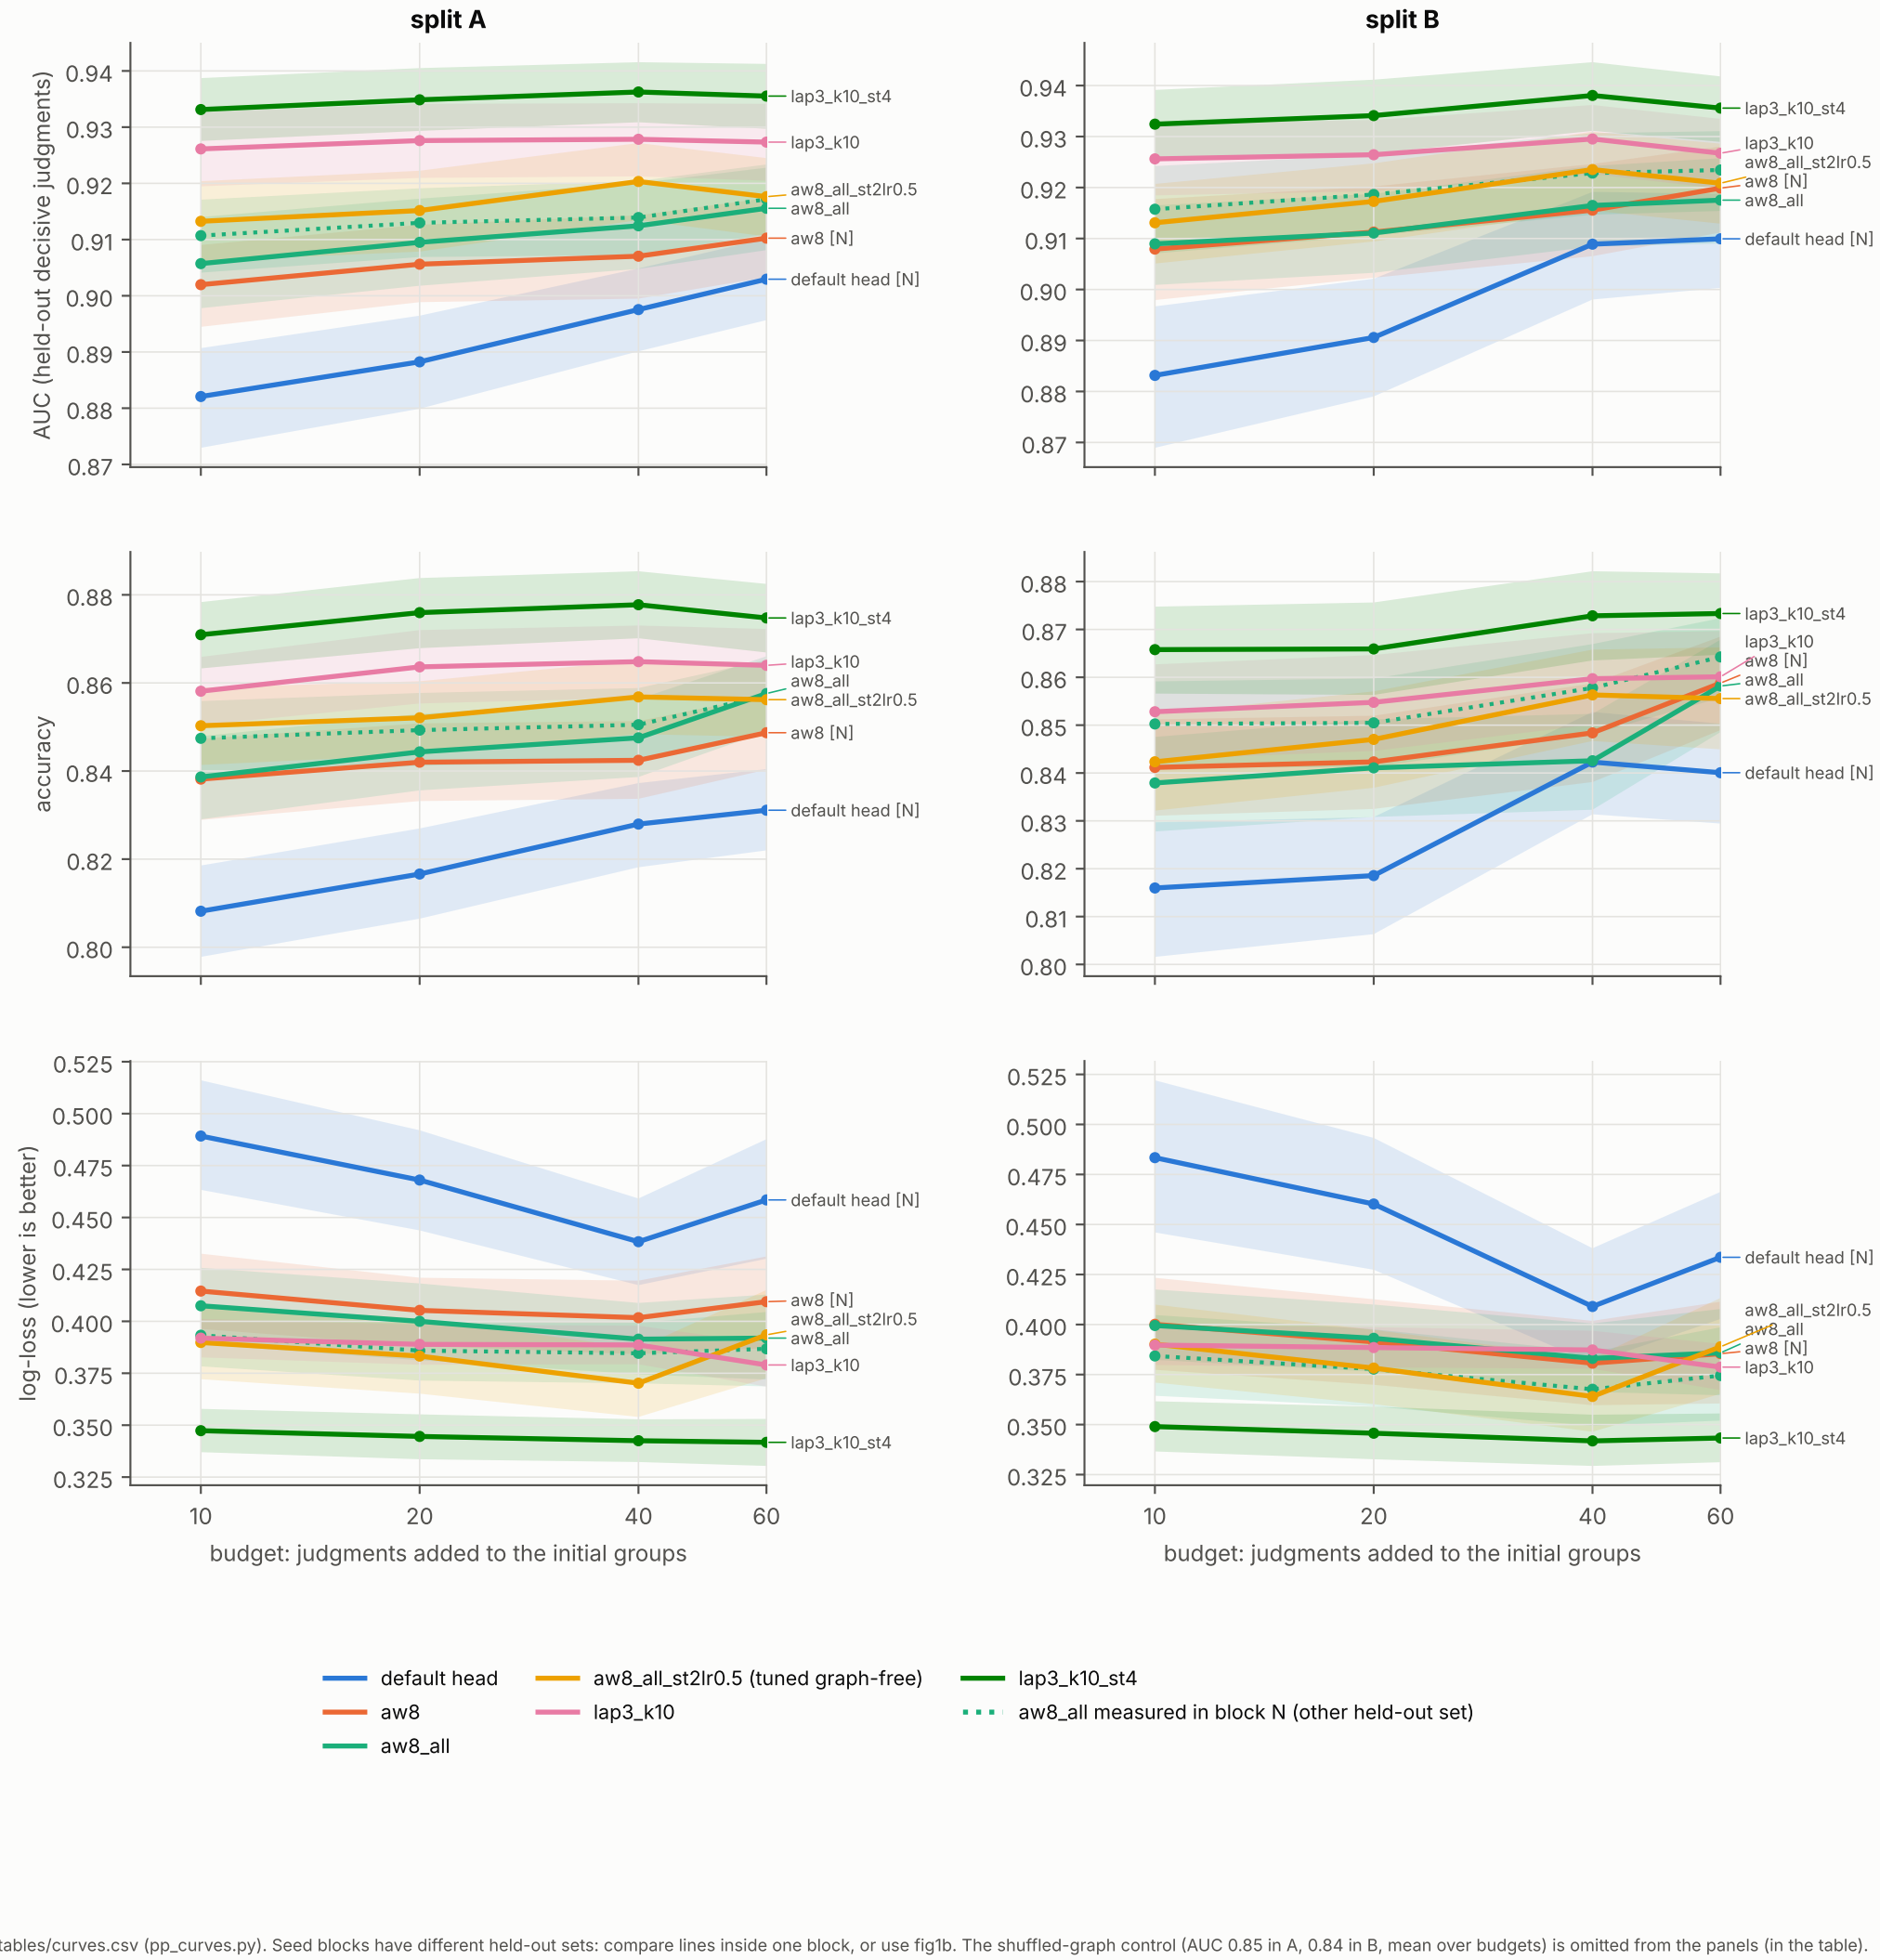

Source: tables/curves.csv (pp_curves.py). Seed blocks have different held-out sets: compare lines inside one block, or use fig1b. The shuffled-graph control (AUC 0.85 in A, 0.84 in B, mean over budgets) is omitted from the panels (in the table).