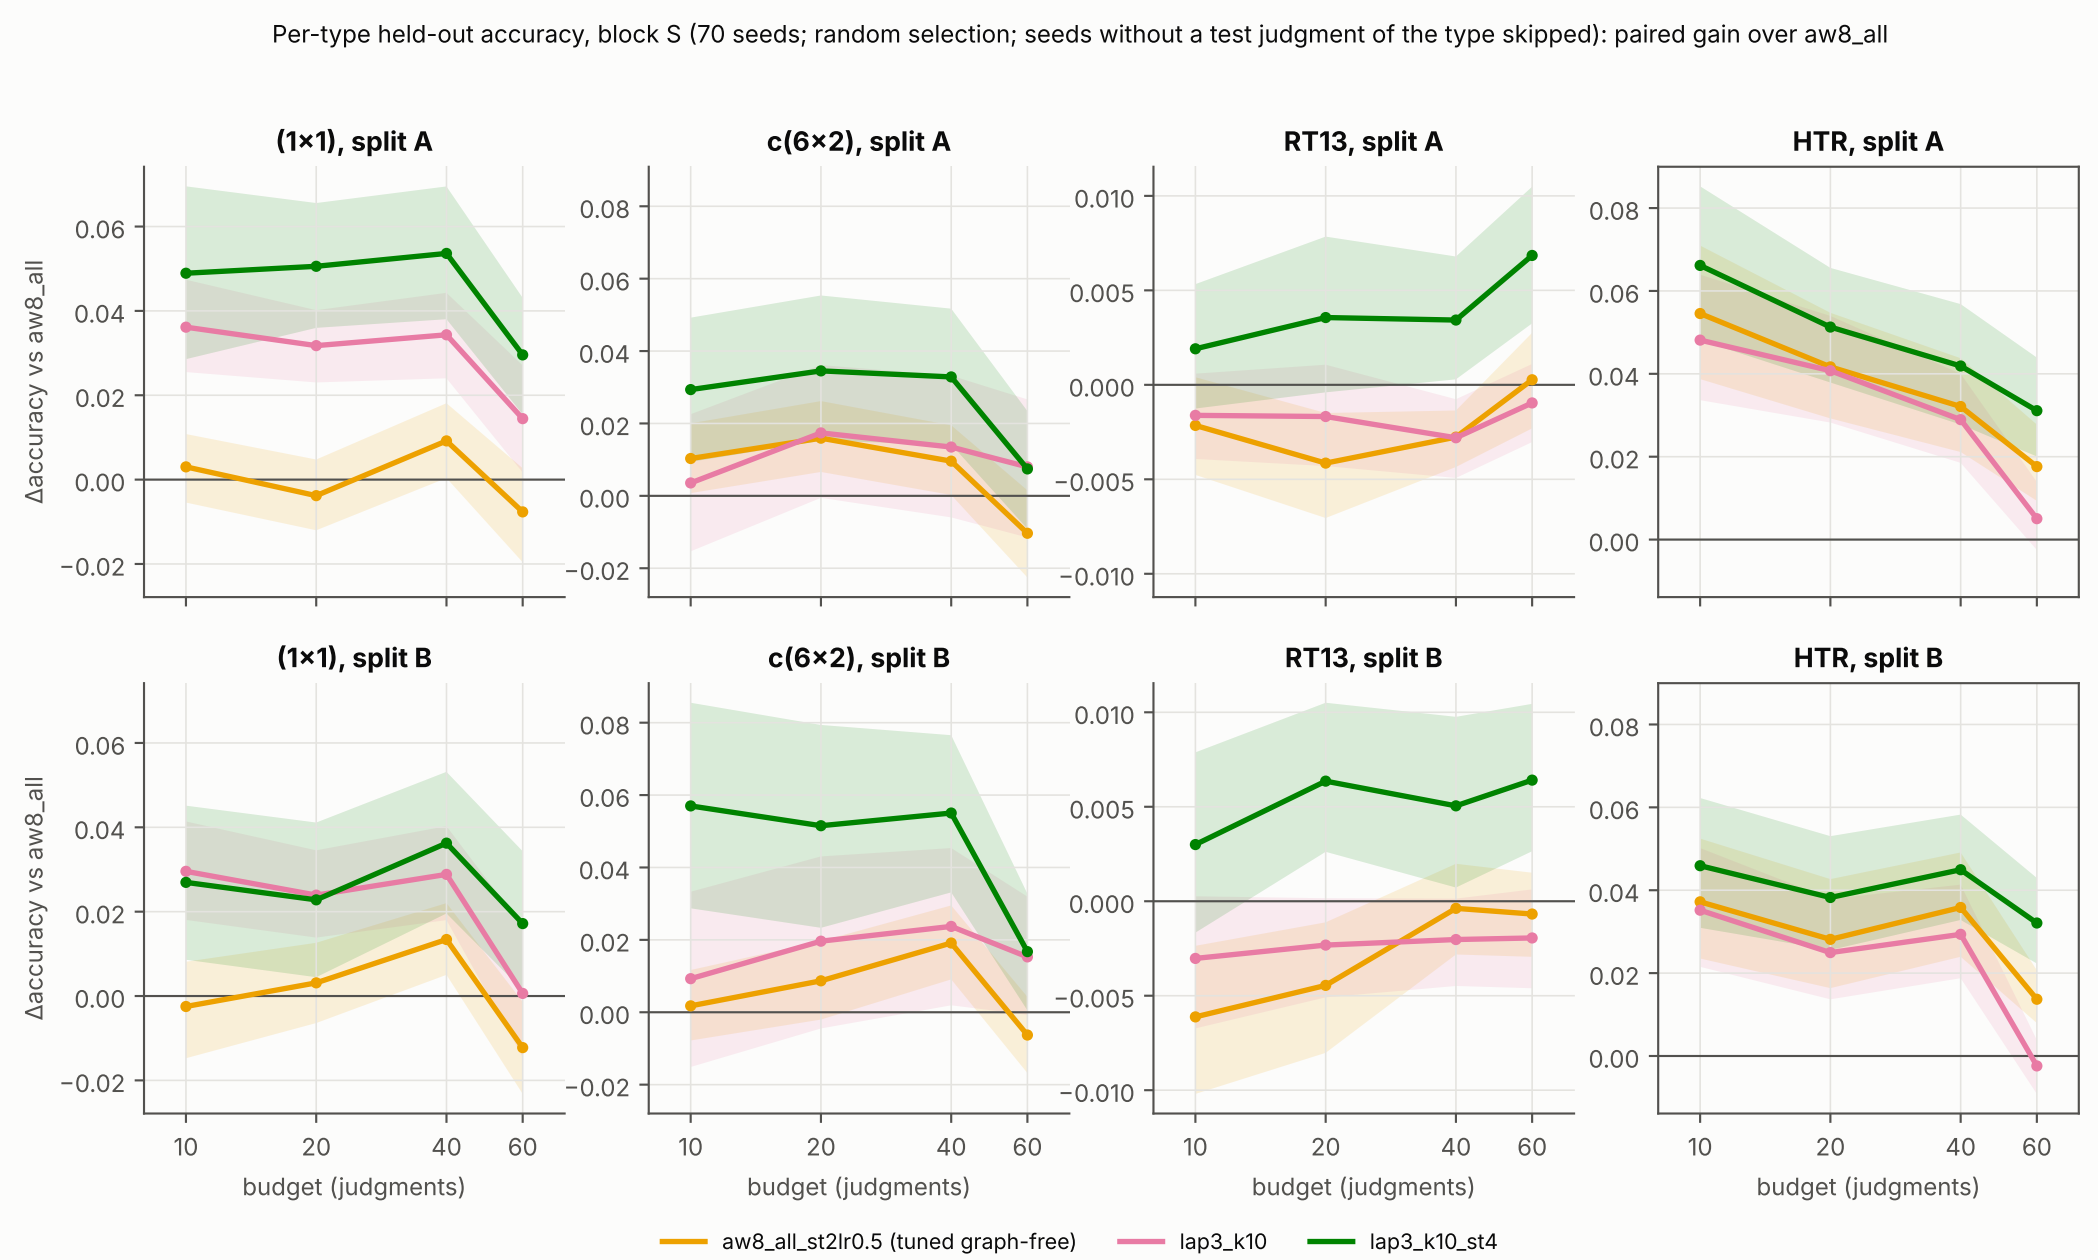
Per-type held-out accuracy, block S (70 seeds; random selection; seeds without a test judgment of the type skipped): paired gain over aw8_all
(1×1), split A
c(6×2), split A
RT13, split A
HTR, split A
(1×1), split B
c(6×2), split B
RT13, split B
HTR, split B
Δaccuracy vs aw8_all
budget (judgments)
aw8_all_st2lr0.5 (tuned graph-free)
lap3_k10
lap3_k10_st4
Source: tables/curves_type.csv (pp_curves.py). HTR panels boxed: the lab's priority type, but only a handful of held-out HTR judgments per seed and split.